

```
1 import pandas as pd
2 from sklearn.preprocessing import OneHotEncoder
3 import seaborn as sb
4 from sklearn.model_selection import train_test_split
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from sklearn.neighbors import KNeighborsClassifier
8 from sklearn.svm import SVC
9 from sklearn.metrics import recall_score, confusion_matrix, precision_score, f1_score, accuracy_score, classification_report
10 from sklearn.linear_model import LogisticRegression
11 from sklearn.neural_network import MLPClassifier
12 from xgboost import XGBClassifier
13 from sklearn.ensemble import RandomForestClassifier
14 from sklearn import metrics
```

```
1 df = pd.read_excel("/content/customer_churn_large_dataset.xlsx")
2 df
```

	CustomerID	Name	Age	Gender	Location	Subscription_Length_Months	Monthly_Bill	To
0	1	Customer_1	63	Male	Los Angeles	17	73.36	
1	2	Customer_2	62	Female	New York	1	48.76	
2	3	Customer_3	24	Female	Los Angeles	5	85.47	
3	4	Customer_4	36	Female	Miami	3	97.94	
4	5	Customer_5	46	Female	Miami	19	58.14	
...
99995	99996	Customer_99996	33	Male	Houston	23	55.13	
99996	99997	Customer_99997	62	Female	New York	19	61.65	
99997	99998	Customer_99998	64	Male	Chicago	17	96.11	
99998	99999	Customer_99999	51	Female	New York	20	49.25	
99999	100000	Customer_100000	27	Female	Los Angeles	19	76.57	

```
1 #Changing into One hot encoding
2 df = pd.get_dummies(df, columns = ['Location', 'Gender'])
3 print(df)
```

	CustomerID	Name	Age	Subscription_Length_Months	\
0	1	Customer_1	63	17	
1	2	Customer_2	62	1	
2	3	Customer_3	24	5	
3	4	Customer_4	36	3	
4	5	Customer_5	46	19	
...
99995	99996	Customer_99996	33	23	
99996	99997	Customer_99997	62	19	
99997	99998	Customer_99998	64	17	
99998	99999	Customer_99999	51	20	
99999	100000	Customer_100000	27	19	

	Monthly_Bill	Total_Usage_GB	Churn	Location_Chicago	\
0	73.36	236	0	0	
1	48.76	172	0	0	
2	85.47	460	0	0	
3	97.94	297	1	0	
4	58.14	266	0	0	
...
99995	55.13	226	1	0	
99996	61.65	351	0	0	
99997	96.11	251	1	1	
99998	49.25	434	1	0	
99999	76.57	173	1	0	

	Location_Houston	Location_Los Angeles	Location_Miami	\
0	0	1	0	
1	0	0	0	
2	0	1	0	
3	0	0	1	
4	0	0	1	
...

```
99995      1      0      0
99996      0      0      0
99997      0      0      0
99998      0      0      0
99999      0      1      0

Location_New York Gender_Female Gender_Male
0      0      0      1
1      1      1      0
2      0      1      0
3      0      1      0
4      0      1      0
...      ...      ...      ...
99995      0      0      1
99996      1      1      0
99997      0      0      1
99998      1      1      0
99999      0      1      0

[100000 rows x 14 columns]
```

```
1 #Finding correlation between all the columns
2 df.corr()

<ipython-input-81-2f6f6606aa2c>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is
df.corr()
```

	CustomerID	Age	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
CustomerID	1.000000	-0.001085	0.005444	0.001265	-0.004025	-0.004586
Age	-0.001085	1.000000	0.003382	0.001110	-0.000666	-0.001390
Subscription_Length_Months	0.005444	0.003382	1.000000	-0.005294	0.002187	-0.001390
Monthly_Bill	0.001265	0.001110	-0.005294	1.000000	0.003187	-0.000211
Total_Usage_GB	-0.004025	0.001927	-0.002203	0.003187	1.000000	-0.000211
Churn	-0.004586	0.001559	0.002328	-0.000211	-0.000211	1.000000
Location_Chicago	-0.000666	0.006068	0.002187	-0.005772	-0.000666	-0.000666
Location_Houston	-0.001390	0.001795	-0.001842	0.001856	-0.001390	-0.001390
Location_Los Angeles	0.002501	-0.004971	-0.001234	0.003444	-0.001234	-0.001234
Location_Miami	0.001617	0.001079	0.005508	-0.002521	0.001617	0.001617
Location_New York	-0.002069	-0.003982	-0.004630	0.002992	0.002992	0.002992
Gender_Female	0.000131	-0.000832	-0.000320	-0.002239	0.000131	0.000131
Gender_Male	-0.000131	0.000832	0.000320	0.002239	-0.000131	-0.000131

```
1 #To give brief about the dataset
2 df.describe()
```

	CustomerID	Age	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
count	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000
mean	50000.500000	44.027020	12.490100	65.053197	274.393650	0.400000
std	28867.657797	15.280283	6.926461	20.230696	130.463063	0.400000
min	1.000000	18.000000	1.000000	30.000000	50.000000	0.000000
25%	25000.750000	31.000000	6.000000	47.540000	161.000000	0.000000
50%	50000.500000	44.000000	12.000000	65.010000	274.000000	0.000000
75%	75000.250000	57.000000	19.000000	82.640000	387.000000	1.000000
max	100000.000000	70.000000	24.000000	100.000000	500.000000	1.000000

```
1 #To find outliers
2 q1 = df['Subscription_Length_Months'].quantile(0.25)
3 q3 = df['Subscription_Length_Months'].quantile(0.75)
4 iqr = q3 - q1
5
6 # Define the lower and upper bounds for outliers
```

```

7 lower_bound = q1 - 1.5 * iqr
8 upper_bound = q3 + 1.5 * iqr
9
10 # Identify the outliers
11 outliers = df[df['Subscription_Length_Months'] < lower_bound] | df[df['Subscription_Length_Months'] > upper_bound]
12
13 # Print the outliers
14 print(outliers)

```

Empty DataFrame

Columns: [Female, Male, Chicago, Houston, Los Angeles, Miami, New York, CustomerID, Name, Age, Subscription_Length_Months, Monthly_Bill, Index: []]

```

1 #Finding outliers in Total_Usage_GB
2 q1 = df['Total_Usage_GB'].quantile(0.25)
3 q3 = df['Total_Usage_GB'].quantile(0.75)
4 iqr = q3 - q1
5
6 # Define the lower and upper bounds for outliers
7 lower_bound = q1 - 1.5 * iqr
8 upper_bound = q3 + 1.5 * iqr
9
10 # Identify the outliers
11 outliers = df[df['Total_Usage_GB'] < lower_bound] | df[df['Total_Usage_GB'] > upper_bound]
12
13 # Print the outliers
14 print(outliers)

```

Empty DataFrame


Columns: [Female, Male, Chicago, Houston, Los Angeles, Miami, New York, CustomerID, Name, Age, Subscription_Length_Months, Monthly_Bill, Index: []]

```

1 sns.pairplot(df, x_vars=['Age', 'Subscription_Length_Months',
2     'Monthly_Bill', 'Total_Usage_GB', 'Location_Chicago',
3     'Location_Houston', 'Location_Los Angeles', 'Location_Miami',
4     'Location_New York', 'Gender_Female', 'Gender_Male'], y_vars='Churn', height=5, aspect=0.7, kind='reg')
5 # There is not much relation between the churn and diiferent columns

```

<seaborn.axisgrid.PairGrid at 0x7b39de90fb80>



```
1 dataplot = sb.heatmap(df.corr(),annot=True )
```

Heatmap visualization showing the correlation matrix for the Telecom Churn dataset. The variables are CustomerID, Age, Subscription_Length_Months, Monthly_Bill, Total_Usage_GB, Churn, and Location_Chicago. The color scale ranges from 0.00 (dark purple) to 1.00 (yellow).

	CustomerID	Age	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn	Location_Chicago
CustomerID	1.000	0.001	0.054	0.130	0.040	0.034	0.000
Age	0.001	1.000	0.008	0.100	0.030	0.160	0.013
Subscription_Length_Months	0.054	0.008	1.000	0.020	0.230	0.220	0.040
Monthly_Bill	0.130	0.100	0.020	1.000	0.010	0.120	0.050
Total_Usage_GB	0.040	0.030	0.230	0.010	1.000	0.160	0.130
Churn	0.034	0.160	0.220	0.120	0.160	1.000	0.020
Location_Chicago	0.000	0.013	0.040	0.050	0.130	0.020	1.000

```
Location_NEW YORK -000.110830030020097005@72072307230723 1 -00233002
lumn_name = 'Churn'
Gender Male -000.0808300300202000.40201000.0062000.40400026-1 1
```

80000 rows x 11 columns

```
1 num_cols = ["Subscription_Length_Months", 'Monthly_Bill', 'Total_Usage_GB']
2 for feat in num_cols: distplot(feat, df)
```

similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(frame[feature], color= color)
```

<ipython-input-93-8c8257b32bab>:4: UserWarning:




`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(frame[feature], color= color)
```



```
1 #Train a KNN classification model with scikit-learn
2 knn_model = KNeighborsClassifier(n_neighbors = 11)
3 knn_model.fit(X_train,y_train)
4 predicted_y = knn_model.predict(X_test)
5 accuracy_knn = knn_model.score(X_test,y_test)
6 print("KNN accuracy:",accuracy_knn)
```

KNN accuracy: 0.5063

```
1 #Train a SVC classification model with scikit-learn
2 svc_model = SVC(random_state = 1)
```

```

3 svc_model.fit(X_train,y_train)
4 predict_y = svc_model.predict(X_test)
5 accuracy_svc = svc_model.score(X_test,y_test)
6 print("SVM accuracy is :",accuracy_svc)

```

```

1 print(classification_report(y_test, predict_y))

```

```

1 #Train a Random Forest classification model with scikit-learn
2 model_rf = RandomForestClassifier(n_estimators=500 , oob_score = True, n_jobs = -1,
3                                 random_state =50,max_leaf_nodes = 30)
4 model_rf.fit(X_train, y_train)
5
6 # Make predictions
7 prediction_test = model_rf.predict(X_test)
8 print (metrics.accuracy_score(y_test, prediction_test))
9

```

```

0.49965

```

```

1 print(classification_report(y_test, prediction_test))

```

```

              precision    recall  f1-score   support

0               0.50         0.68         0.58         10080
1               0.49         0.32         0.39          9920

 accuracy               0.50         0.50         0.48         20000
 macro avg              0.50         0.50         0.48         20000
 weighted avg           0.50         0.50         0.48         20000

```

```

1 #Train a XGB classification model with scikit-learn
2 model = XGBClassifier()
3 model.fit(X_train, y_train)
4 preds = model.predict(X_test)
5 metrics.accuracy_score(y_test, preds)

```

```

0.50445

```

```

1 #Fine tuning XB Classification model to get better accuracy score
2 model = XGBClassifier(learning_rate=0.0001,
3                       colsample_bytree = 0.4,
4                       subsample = 0.8,
5                       objective='binary:logistic',
6                       n_estimators=1000,
7                       reg_alpha = 0.3,
8                       max_depth=5,
9                       gamma=100)
10 model.fit(X_train, y_train)
11 preds = model.predict(X_test)
12 metrics.accuracy_score(y_test, preds)

```

```

0.504

```

```

1 #Train a MLP classification model with scikit-learn
2 mlp_classifier = MLPClassifier(hidden_layer_sizes=(8, 20), activation='relu', solver='adam', max_iter=50000)
3 mlp_classifier.set_params(alpha=0.0001, batch_size=10, learning_rate_init=0.0001)
4 mlp_classifier.fit(X_train, y_train)
5 y_pred = mlp_classifier.predict(X_test)
6 accuracy = mlp_classifier.score(X_test, y_test)
7
8 print('Accuracy:', accuracy)

```

```

Accuracy: 0.50395

```

```

1 #Train a Logistic Regression model with scikit-learn
2 lr_model = LogisticRegression()
3 lr_model.fit(X_train,y_train)
4 accuracy_lr = lr_model.score(X_test,y_test)
5 print("Logistic Regression accuracy is :",accuracy_lr)

```

```

Logistic Regression accuracy is : 0.5021

```

```

1 lr_pred = lr_model.predict(X_test)

```

```
1 lr_pred= lr_model.predict(X_test)
2 report = classification_report(y_test,lr_pred)
3 print(report)
```

	precision	recall	f1-score	support
0	0.51	0.60	0.55	10080
1	0.50	0.40	0.45	9920
accuracy			0.50	20000
macro avg	0.50	0.50	0.50	20000
weighted avg	0.50	0.50	0.50	20000

```
1 #To improve the accuracy we can use Pytorch Neural Network and can overfit the data ,
2 #but that would be good to only improve the accuracy score but it will give inaccurate result for new data.
3 #Out of all the models implemented KNN Classification works best for the data as it gives highest Accuracy score of 50.63%
```